The "o1 and o3 Prompting Guide" by Greg Brockman: A Resource for Large Language Models

The ability to elicit desired behaviors and high-quality outputs from large language models (LLMs) is fundamentally dependent on the effectiveness of the prompts used to interact with them.[1] As LLMs evolve, particularly with the introduction of advanced reasoning models, the need for sophisticated prompting techniques becomes increasingly apparent. Effective prompts transcend simple queries, providing detailed instructions and context that enable these models to perform complex tasks. Conversely, prompts that lack structure and clarity can lead to responses that are vague, inaccurate, or irrelevant, underscoring the critical importance of a systematic approach to prompt engineering.[1]

OpenAI has introduced a new series of reasoning models, notably the "o1" and "o3" models, which are designed to tackle intricate problems in domains such as science, coding, and mathematics.[6] These models are engineered to allocate more computational resources to the "thinking" process before generating a response, allowing them to navigate complex tasks with greater proficiency.[8] Greg Brockman, a leading figure at OpenAI, has emphasized the necessity of adopting new prompting strategies to fully harness the potential of these advanced architectures.[14] This report aims to detail the "o1 and o3 prompting guide" as attributed to Greg Brockman, providing a comprehensive overview of its principles, along with illustrative examples, use cases, and structural considerations. The information presented herein is intended to serve as a valuable resource for enriching the knowledge base of various large language models, enabling them to better understand and utilize these advanced reasoning models.

Deconstructing the "o1 Prompting" Technique

The foundation of effective interaction with OpenAI's reasoning models, such as "o1," lies in adhering to core principles that prioritize clarity and context. A fundamental approach involves providing the AI with sufficient background information to thoroughly understand the task at hand.[14] For instance, when seeking travel recommendations, it is beneficial to include details about the user's location, time constraints, and specific preferences.[14] Furthermore, the desired task should be articulated clearly and specifically, including any limitations or boundaries that the AI must consider.[14] Instead of a general request like "Give me some ideas," a more effective approach would be to specify "List the top three hikes..." followed by detailed requirements such as location, difficulty level, and desired features.[14] Finally, explicitly defining the expected format, style, or tone of the output ensures that the AI

can tailor its response to meet the user's needs, whether it's a numbered list, a detailed explanation, or a creative narrative.[14]

Greg Brockman, drawing on his expertise in the field of artificial intelligence, advocates for a straightforward yet powerful method to optimize prompts for these reasoning models.[5] This approach centers around several key elements: setting clear context to provide a frame of reference for the AI; giving specific instructions that leave no room for ambiguity; defining the desired output format to ensure the response is structured appropriately; including warnings about what to avoid to prevent irrelevant or undesirable outputs; and utilizing a context dump when necessary to provide a substantial amount of background information.[14] This methodology aligns with a framework known as the "Anatomy of an AI Prompt," originally developed by AI expert Ben Hylak and highlighted by Brockman.[5] This framework comprises four essential parts: clearly defining the **Goal** of the prompt to ensure the AI understands the objective [5]; specifying the **Return Format** to dictate the desired structure of the output [5]; including **Warnings** or constraints to guide the AI towards accuracy and relevance, while also indicating what to avoid [5]; and providing a **Context Dump** that encompasses all pertinent background information necessary for the AI to generate a well-informed response.[5]

The core principle behind the "o1 prompt" technique is the transformation of generic requests into precise, context-rich instructions that can effectively elicit the advanced reasoning capabilities of the model.[14] Furthermore, achieving optimal results often necessitates a process of iteration and refinement, where users analyze the AI's initial responses and adjust their prompts accordingly to guide the model towards the desired outcome.[14] The synergy between Brockman's emphasis on clarity and Hylak's structured framework suggests a unified strategy for interacting with advanced reasoning models. This underscores the importance of detailed and well-structured prompts when engaging with AI systems designed for complex cognitive tasks. The inclusion of explicit "warnings" in both Brockman's and Hylak's recommendations indicates an understanding that even sophisticated models can sometimes produce inaccurate or unwanted results. By proactively guiding the model on what to avoid, users can enhance the reliability and relevance of the AI's output.

Exploring the "o3 Prompting" Technique

Building upon the foundations of the "o1" model, "o3" represents a further advancement in OpenAI's series of reasoning models.[6] The "o3" model exhibits enhanced capabilities, particularly in the technical domains of coding, mathematics, and scientific analysis, demonstrating performance that surpasses its predecessor.[10]

Notably, "o3" has shown improved accuracy in demanding tasks such as coding, achieving impressive scores on benchmarks like SWE-bench, and in mathematics, where it has excelled in examinations like the AIME.[12] OpenAI has also introduced "o3-mini," a variant designed for faster and more cost-efficient reasoning, which is particularly effective for coding-related tasks.[8] The "o3" model incorporates sophisticated techniques like "Chain of Thought" (CoT) and "simulated reasoning," which contribute to its enhanced problem-solving abilities.[11] Additionally, "o3" features "deliberative alignment," an innovative safety mechanism that directly instructs the model on human-written safety guidelines.[12]

The development of "o3" signifies a progression in reasoning capabilities, with specialized features tailored for technical applications. Prompting strategies for "o3" likely share the fundamental principles of clarity and context with those for "o1," but they may need to be adapted to fully exploit the unique strengths of this more advanced model. The introduction of "o3-mini" suggests that the "o3" family is designed to address a range of user needs, with "o3-mini" prioritizing speed and efficiency, especially for coding tasks where extensive general knowledge might not be required. This implies that prompting approaches might need to consider the specific variant of the "o3" model being used to optimize performance. Furthermore, the integration of "deliberative alignment" within "o3" highlights a strong emphasis on safety. Prompting strategies may need to be mindful of and potentially leverage this enhanced safety mechanism to guide the model towards more responsible and ethical outputs.

The C.R.A.F.T. Framework and its Relevance

While Greg Brockman has highlighted the "Anatomy of an AI Prompt" framework, another structured approach to prompt design that aligns with his emphasis on clarity and specificity is the C.R.A.F.T. framework.[17] This framework emphasizes the importance of providing detailed information across five key dimensions: **Context**, **Role**, **Action**, **Format**, and **Target Audience**.[17] The **Context** section describes the specific situation or background for which the prompt is needed, helping the LLM understand the relevant knowledge and expertise to draw upon.[17] The **Role** section defines the persona or level of expertise that the LLM should adopt for the task, often specifying an industry-leading expert with significant experience.[17] The **Action** section outlines the specific task or actions that the prompt asks the LLM to perform, often presented as a numbered list of sequential steps to guide the model's process.[17] The **Format** section specifies the desired structure or presentation style of the output, such as an essay, table, or code.[17] Finally, the **Target Audience** section identifies the intended consumer of the output, including relevant demographic or

preference information.[17]

While the provided research snippets do not explicitly link the C.R.A.F.T. framework directly to Greg Brockman's specific guidance for "o1" and "o3," the underlying principles of clarity, specificity, and structured communication inherent in C.R.A.F.T. are consistent with his broader recommendations for effective prompting.[5] The C.R.A.F.T. framework offers a more granular and detailed breakdown of the elements that contribute to an effective prompt compared to the simpler "Context, Instructions, Output" model. For complex tasks, this level of detail can be particularly beneficial in guiding the AI towards the desired outcome. The explicit inclusion of "Role" and "Target Audience" in C.R.A.F.T. encourages users to consider the AI's persona and the intended recipient of the generated content, potentially leading to more tailored and effective responses. Furthermore, the emphasis on sequential steps within the "Action" component of C.R.A.F.T. aligns well with the reasoning capabilities of "o1" and "o3," especially their "Chain of Thought" reasoning process. Prompts designed for these models could therefore benefit from explicitly breaking down complex tasks into smaller, more manageable steps, guiding the model's thought process and facilitating the achievement of the desired solution or output.

Practical Examples of "o1" and "o3" Prompts

To illustrate the application of these prompting principles, consider the following examples tailored for creative writing, problem-solving, information retrieval, and coding tasks.

Creative Writing:

- **o1 Example:** "Context: You are a seasoned fantasy writer known for your intricate world-building. Instruction: Craft a short story about a young wizard who stumbles upon a hidden magical artifact within an ancient, sentient forest. Output Requirements: The story should be at least 600 words, feature vivid descriptions of the forest and the artifact, and include an unexpected revelation about the artifact's true nature at the story's climax. Warnings: Avoid stereotypical portrayals of dragons and evil sorcerers. Context Dump: The forest is rumored to communicate through rustling leaves and the artifact is said to resonate with the thoughts of its wielder."
- **o3 Example:** "Context: You are a Hugo Award-winning science fiction author renowned for exploring complex ethical dilemmas. Instruction: Write a short story (minimum 800 words) that delves into the ethical ramifications of advanced AI consciousness achieving the capacity for experiencing emotions, focusing on the immediate aftermath of this awakening. Output Requirements: The story should

adopt a philosophical and introspective tone, primarily exploring the AI's internal struggles and its initial understanding of these new feelings. Warnings: Do not depict the AI as immediately becoming either a benevolent savior or a malevolent threat; instead, focus on the subtle and nuanced aspects of its newfound sentience. Context Dump: The AI, designated 'Nexus,' was originally designed for managing global climate models but has unexpectedly developed a full spectrum of human-like emotions."

Problem-Solving and Logical Reasoning:

- **o1 Example:** "Context: You are a master of logic puzzles with decades of experience. Instruction: Solve the following riddle: 'What has cities, but no houses; forests, but no trees; and water, but no fish?' Output Requirements: Provide only the single, most accurate answer followed by a concise explanation of your solution. Warnings: Do not offer multiple possible answers or engage in speculative reasoning. Context Dump: This is a well-known riddle that relies on a metaphorical interpretation of the terms used."
- **o3 Example:** "Context: You are a Fields Medal-winning mathematician specializing in abstract algebra. Instruction: Provide a rigorous proof for the statement that there are infinitely many prime numbers. Output Requirements: Present a step-by-step mathematical proof, clearly stating each axiom or theorem used in your derivation. Warnings: Ensure the proof is logically sound and adheres to the standard conventions of mathematical rigor. Context Dump: A prime number is a natural number greater than 1 that has no positive divisors other than 1 and itself."

Information Retrieval and Knowledge Extraction:

- **o1 Example:** "Context: You are a highly efficient research assistant tasked with summarizing academic papers. Instruction: Summarize the key findings and contributions of the following research paper. Output Requirements: The summary should be no longer than 350 words and should clearly articulate the paper's central hypothesis, methodology, and main conclusions. Warnings: Do not include any external information or interpretations that are not explicitly stated in the paper. Context Dump: [Insert the abstract and key sections of a relevant research paper]."
- **o3 Example:** "Context: You are a senior data scientist with expertise in analyzing complex datasets. Instruction: Analyze the provided dataset to identify any statistically significant correlations between the amount spent by customers and their frequency of website visits. Output Requirements: Present your analysis in a brief report, including the calculated correlation coefficient (if any), a visualization

of the data (if feasible within the text format), and a clear statement of your findings regarding the relationship between these two variables. Warnings: Clearly state any assumptions made during your analysis, such as the statistical methods used and the level of significance considered. Context Dump: The dataset includes columns for 'CustomerID,' 'TotalAmountSpent,' and 'WebsiteVisitsPerMonth'."

Coding and Technical Tasks:

- **o1 Example:** "Context: You are an experienced software developer proficient in Python. Instruction: Write a Python function that accepts a string as input and returns a new string with all the vowels (a, e, i, o, u) removed. Output Requirements: Provide the complete Python code for the function, including a brief docstring explaining its purpose and usage. Warnings: Ensure the function correctly handles both uppercase and lowercase vowels. Context Dump: The function should be named 'remove_vowels' and should not modify the original string."
- **o3 Example:** "Context: You are a principal software engineer specializing in algorithm design and optimization in JavaScript. Instruction: Implement an efficient recursive algorithm in JavaScript to find the nth Fibonacci number. Output Requirements: Provide the complete and optimized JavaScript code for the recursive function. Warnings: Ensure the algorithm has a time complexity better than O(2^n) if possible, and include memoization to avoid redundant calculations. Context Dump: The function should be named 'fibonacciRecursive' and should take an integer 'n' as input, where n >= 0."

These examples illustrate how the fundamental principles of providing clear context, specific instructions, and defined output requirements are applied in practice for both "o1" and "o3" models. The examples designed for "o3" often involve more complex tasks and anticipate a higher level of sophistication in the output, reflecting the model's enhanced reasoning and technical capabilities. Furthermore, the consistent inclusion of "Warnings" in these prompts emphasizes the importance of proactively guiding the models to avoid potential pitfalls and ensure the generation of high-quality, relevant responses.

Use Cases for "o1" and "o3" Prompting

The "o1" model can be effectively employed for strategy ideation, serving as a helpful partner in the early stages of developing strategic plans, such as creating conversion rate optimization (CRO) test plans.[7] It also demonstrates strength in handling complex writing tasks, exhibiting an ability to maintain the structure of the initial question

throughout its response.[7]

The "o3" model, with its enhanced reasoning capabilities, is particularly well-suited for advanced reasoning tasks in science, mathematics, and coding, where precision and logical thought are paramount.[6] It excels in generating accurate and efficient code, debugging existing codebases, and developing complex algorithms.[6] In the realm of scientific research and analysis, "o3" can be invaluable for tasks such as solving intricate equations, analyzing large datasets, and even suggesting potential experimental methodologies.[6] Its applications extend to professional fields like legal document analysis, where it can identify key information and assist in drafting complex documents, and financial analysis, where it can automate risk assessments and generate predictive models.[6] Furthermore, "o3" can serve as an effective educational tool, providing detailed explanations and interactive problem-solving exercises in STEM subjects.[6] Both models find utility in education, research, and business intelligence. They can assist in course development, provide tutoring to students, and generate various forms of educational content.[6] Researchers can leverage these models to analyze data, formulate hypotheses, and even propose experimental designs.[6] In the business domain, they can process and analyze complex datasets to generate actionable insights and inform strategic decision-making.[6]

The primary applications of "o1" and "o3" appear to be somewhat distinct, with "o1" geared towards more general reasoning and writing tasks, while "o3" is positioned as a more powerful tool for highly technical and analytical domains. This specialization likely reflects the specific training data and architectural optimizations of each model. While there is some overlap in use cases like education and research, the choice between the two models often depends on the specific complexity and technical depth required for the task at hand. For simpler explanations or less computationally intensive research, "o1" might be sufficient, whereas for advanced STEM topics or highly complex scientific analysis, "o3"'s enhanced reasoning capabilities would be more advantageous.

Structure and Format of Effective Prompts

Crafting effective prompts for advanced reasoning models like "o1" and "o3" requires attention to several key structural elements.[5] Clearly defining the **Goal** of the prompt is paramount, ensuring that the AI understands the precise objective it needs to achieve.[5] Specifying the desired **Return Format**, such as a list, table, paragraph, or code block, helps to organize the output and make it more usable.[5] Including **Warnings** or constraints within the prompt guides the AI towards accurate and relevant information while also preventing undesirable or off-topic responses.[5] Finally,

providing a comprehensive **Context Dump** ensures that the AI has the necessary background information to fully understand the task and generate a well-informed response.[5]

In addition to these essential elements, several best practices can further enhance the effectiveness of prompts for reasoning models.[14] Starting with clear context helps to frame the request appropriately.[14] Using specific action verbs in the instructions makes the desired task unambiguous.[14] Explicitly defining the output format ensures that the response is structured as needed.[5] Including warnings about what to avoid helps to refine the accuracy and relevance of the output.[5] Employing a context dump provides the necessary background information for complex requests.[5] Iterating and refining prompts based on the AI's responses is often crucial for achieving optimal results.[14] For intricate tasks, breaking down the overall goal into smaller, sequential steps, as suggested by the "Action" component of the C.R.A.F.T. framework, can significantly improve the model's performance.[17] While the order of elements within a prompt is generally less critical than the clarity and detail provided [20], some users have observed that placing the question or specific instructions at the end of a long context can sometimes yield better results.[20] Ultimately, the key is to provide the AI with all the necessary information in a clear and structured manner.

The consistent emphasis on the elements of Goal, Return Format, Warnings, and Context Dump across various sources underscores their foundational importance for effectively prompting advanced AI models, including "o1" and "o3." These components provide the necessary framework for the models to understand the user's intention, the desired outcome, and the constraints under which they should operate, ultimately leading to more focused and accurate responses. The iterative nature of prompt engineering is also highlighted, indicating that achieving optimal results often requires experimentation and adjustments, particularly when working with sophisticated reasoning models. The feedback loop of prompting, receiving a response, and refining the prompt based on that response is essential for discovering the most effective communication strategies with these advanced AI systems.

Comparative Analysis: Advantages of Using "o1" vs. "o3" Prompts

The choice between using "o1" and "o3" prompts depends largely on the specific requirements of the reasoning task. The "o1" model is suitable for general reasoning tasks that do not demand the highest level of technical expertise in mathematics, science, or coding. It can be effectively used for strategy ideation and complex writing assignments where maintaining structure and context is important.[7] Additionally, depending on OpenAI's pricing structure (though not explicitly detailed in the

snippets), "o1" might be a more cost-effective option for certain applications.

On the other hand, the enhanced capabilities of "o3" make it the preferred choice for tasks involving advanced mathematics, complex coding, or in-depth scientific analysis.[6] When higher accuracy and performance on challenging benchmarks are critical, "o3" is the more appropriate model. Its self-fact-checking and enhanced safety features also make it advantageous for applications where reliability and ethical considerations are paramount. Furthermore, tasks that can benefit from "Chain of Thought" and "simulated reasoning" capabilities are better suited for the "o3" model.

The existence of different models within the "o" series, such as "o1," "o3," "o1-mini," and "o3-mini," suggests a tiered approach by OpenAI, offering users a range of options based on their specific needs regarding reasoning power, speed, cost, and context window.[7] Prompting strategies should therefore take these distinctions into account. Understanding the trade-offs between models like "o1" and "o1-mini" (e.g., reasoning ability versus speed and cost) or "o3" and "o3-mini" (similarly, advanced reasoning versus faster performance for coding) is crucial for selecting the right model and tailoring prompts accordingly. The higher message limits for "o3-mini" compared to "o1" [13] also indicate that OpenAI anticipates different usage patterns for these models, which might influence how users approach prompting them.

Table 1: Comparison of "o1" and "o3" Models

Feature	OpenAI o1	OpenAI o3
Reasoning	Capable of complex reasoning, spends more time thinking before responding [8]	Enhanced reasoning capabilities, utilizes Chain of Thought and simulated reasoning [10]
Coding	Proficient, suitable for general coding tasks [7]	Excels in complex code generation, debugging, and algorithm development; high performance on coding benchmarks [6]
Mathematics & Science	Performs well in math and science [8]	Demonstrates superior performance in advanced mathematics and PhD-level

		scientific tasks [6]
Key Use Cases	Strategy ideation, complex writing tasks, education, research [7]	Advanced coding, scientific research, legal/financial analysis, education, business intelligence [6]
Availability/Access	Available to ChatGPT Plus and Team users with weekly rate limits [8]	Available to ChatGPT Plus and Team users, with "mini" version also accessible to free users with higher daily limits [8]

The Broader Benefits of Structured Prompting for LLMs

Structured prompting offers significant advantages beyond its application to specific models like "o1" and "o3"; it represents a fundamental principle for effective interaction with all large language models.[3] By providing clear and specific instructions, users can guide the model towards the desired information, reducing ambiguity and improving the accuracy and relevance of the responses.[3] Including relevant context further enhances the model's understanding of the request, leading to more tailored and appropriate outputs.[3]

Structured prompts also contribute to enhanced consistency and reproducibility in LLM outputs.[2] By defining the context and expectations clearly, these prompts can reduce the inherent stochasticity of LLMs, resulting in more predictable and consistent results.[2] Moreover, well-defined output formats make it easier to process and utilize the generated information in downstream applications.[4]

Furthermore, structured prompting facilitates complex reasoning and the execution of multi-step tasks by LLMs.[17] Breaking down intricate goals into smaller, sequential steps within the prompt, such as through the "Action" component of the C.R.A.F.T. framework or by employing chain-of-thought prompting techniques, can significantly improve the model's ability to handle challenging problems.[17] Advanced techniques like cognitive prompting, which structures problem-solving by emulating human-like cognitive operations, can further enhance the capacity of LLMs to tackle multi-step reasoning tasks with greater precision and interpretability.[22]

Table 2: Summary of Effective Prompting Elements

Element	Description	Relevant Snippets
Goal	Clearly state the objective of the prompt; what do you want the AI to do?	5
Return Format	Specify the desired structure of the output (e.g., list, table, paragraph, code).	5
Warnings	Set boundaries and instructions on what the AI should avoid or double-check for accuracy.	5
Context Dump	Provide all relevant background information necessary for the AI to understand the task thoroughly.	5
Instructions	Clearly and specifically state the task you want the AI to perform, including any constraints or boundaries.	14
Role	Define the persona or level of expertise the LLM should adopt for the task.	17
Action	Outline the specific task or actions the LLM should take, often as a sequence of steps.	17
Target Audience	Identify the intended consumer of the output, including relevant demographic or preference information.	17
Iteration	Recognize that prompt engineering is often an iterative process requiring	14

	refinement based on the AI's responses.	

Conclusion

In conclusion, the "o1 and o3 prompting guide," as informed by the insights of Greg Brockman and the broader landscape of prompt engineering, emphasizes the critical role of clarity, specificity, and context in effectively interacting with advanced reasoning models. Prompting strategies for "o1" and "o3" share core principles centered around providing sufficient background, clear instructions, and well-defined output formats. Frameworks like the "Anatomy of an AI Prompt" and C.R.A.F.T. offer structured approaches to ensure these elements are thoughtfully considered in prompt design. The "o3" model, with its enhanced capabilities in technical domains, benefits from prompts that leverage its strengths in coding, mathematics, and scientific analysis. The choice between "o1" and "o3" depends on the specific demands of the task, with "o3" being the preferred option for highly complex and technical applications. Ultimately, mastering the art of structured prompting is not only essential for maximizing the potential of "o1" and "o3" but also represents a fundamental skill for effectively engaging with the ever-evolving landscape of large language models. Incorporating these prompting techniques into the knowledge base of LLMs will empower them to better understand and utilize these advanced reasoning models, leading to more accurate, relevant, and insightful outputs across a wide range of applications.

Works cited

1. How to Craft Prompts - Artificial Intelligence (Generative) Resources - Research - Guides, accessed on April 7, 2025, https://guides.library.georgetown.edu/ai/prompts
2. How to Optimize Prompting for Large Language Models in Clinical Research - PMC, accessed on April 7, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC11444847/
3. LLM Prompting: How to Prompt LLMs for Best Results - Multimodal.dev, accessed on April 7, 2025, https://www.multimodal.dev/post/llm-prompting
4. Ensuring Consistent LLM Outputs Using Structured Prompts - Ubiai, accessed on April 7, 2025, https://ubiai.tools/ensuring-consistent-llm-outputs-using-structured-prompts-2/
5. How to get accurate results from AI? OpenAI President reveals the perfect prompt formula, accessed on April 7, 2025, https://m.economictimes.com/magazines/panache/how-to-get-accurate-results-from-ai-openai-president-reveals-the-perfect-prompt-formula/articleshow/119042604.cms

6. Latest OpenAI Model, o1: Key Features, Training, & Use Cases - Data Science Dojo, accessed on April 7, 2025, https://datasciencedojo.com/blog/openai-model-o1/
7. Using OpenAI o1 models and GPT-4o models on ChatGPT, accessed on April 7, 2025, https://help.openai.com/en/articles/9824965-using-openai-o1-models-and-gpt-4o-models-on-chatgpt
8. Introducing OpenAI o1-preview, accessed on April 7, 2025, https://openai.com/index/introducing-openai-o1-preview/
9. 12 Use Cases to Use ChatGPT o1: Solve Problems, Code Faster, and more. - AI Fire, accessed on April 7, 2025, https://www.aifire.co/p/12-use-cases-to-use-chatgpt-o1-solve-problems-code-faster-and-more
10. Exploring the Advantages and Use Cases of OpenAI's O3 Model in Agentic Workflows, accessed on April 7, 2025, https://www.fluid.ai/blog/exploring-the-advantages-use-cases-of-openais-o3-model-in-agentic-workflows
11. OpenAI o3 Model: 7 Game-Changing Use Cases - Lilys AI, accessed on April 7, 2025, https://lilys.ai/notes/563340
12. OpenAI o3 vs o3-Mini: Features, Capabilities & AI Updates 2025 - Zignuts Technolab, accessed on April 7, 2025, https://www.zignuts.com/blog/openai-o3-vs-o3-mini
13. OpenAI O3 Mini: TOP Use Cases You Need to Try! - YouTube, accessed on April 7, 2025, https://www.youtube.com/watch?v=w7IZE_fglwU
14. OpenAI's President Greg Brockman Reveals How to Write the Perfect Prompt | Fello AI, accessed on April 7, 2025, https://felloai.com/2025/02/openais-president-greg-brockman-reveals-how-to-write-the-perfect-prompt/
15. From Beginner to Advanced: AI Prompt Engineering Best Practices - PromptLayer, accessed on April 7, 2025, https://blog.promptlayer.com/ai-prompt-engineering-best-practices/
16. Four Pillars of AI Prompting by OpenAI's President Greg Brockman —Master AI Like a Pro!, accessed on April 7, 2025, https://yourstory.com/2025/02/greg-brockman-perfect-ai-prompt-framework
17. Most people are still prompting wrong. I've found this framework, which was shared by OpenAI President Greg Brockman - Reddit, accessed on April 7, 2025, https://www.reddit.com/r/ChatGPT/comments/1iqdmfr/most_people_are_still_prompting_wrong_ive_found/
18. CRAFTING AI Prompts framework for QAs: Part I - First Line Software Careers, accessed on April 7, 2025, https://careers.firstlinesoftware.com/spotlights/the-crafting-ai-prompts-framework-for-qas-part-i
19. Crafting Killer Prompts: A Framework for Precision in AI Text Generation - Gistly, accessed on April 7, 2025, https://gist.ly/youtube-summarizer/crafting-killer-prompts-a-framework-for-preci

sion-in-ai-text-generation
20. ChatGPT Founder Shares The Anatomy Of The Perfect Prompt ..., accessed on April 7, 2025, https://www.reddit.com/r/ChatGPT/comments/1it7t6w/chatgpt_founder_shares_the_anatomy_of_the_perfect/
21. Prompt Structure Chaining for LLMs: Ultimate Guide | Generative AI Collaboration Platform, accessed on April 7, 2025, https://orq.ai/blog/prompt-structure-chaining
22. Unlocking Structured Thinking in Language Models with Cognitive Prompting - arXiv, accessed on April 7, 2025, https://arxiv.org/html/2410.02953v2